

High Performance Computer Architecture (CS60003)

Assignment - 1

Group - 1

Indian Institute of Technology
Kharagpur

GROUP MEMBERS & THEIR CONTRIBUTIONS

Name	Roll Number	Contribution
Utkarsh Patel	18EC35034	Analyzing the results and writing code for plotting.
Aaditya Agrawal	19CS10003	Writing the code to plot the results (analysis.py), analyzing the top 10 configurations and stats.txt and preparing the report.
Akshat Vijay	19CS10008	Analyzing the statistics.
Anindya Sikdar	19CS10010	Writing the custom configuration code(config.py), automating the execution of all parameter combinations(run.py) and preparing the report.
Chandra Pavan Sai Chowdary	19CS10023	Analyzing the top 10 configurations and benchmark program, and preparing the report.
Gandhi Abhishek Rajesh	19CS10031	Modifying options.py and analyzing the results.
Gawai Laukik Thageshwar	19CS10032	Analyzing the top 10 configurations and preparing the report.
Kotaru Parasurama Sai Mani Surya	19CS10041	Automating the execution of parameter combinations and preparing the report.

Note: Each member simulated 32 configuration combinations in their system as each run of configuration was computationally expensive and took a lot of time.

OBJECTIVE

The gem5 simulator is an open-source system-level and processor simulator

supporting multiple ISAs like ARM, MIPS and RISC-V. In this assignment, we configured an out of order CPU with a list of different microarchitectural parameters combinations and ran the benchmarked program on it. The different parameter combinations are automated through a script and the results are sorted based on the CPI. Certain performance parameters are extracted from the generated stats file of top 10 configurations and are plotted to observe their variations.

CODE SUBMISSION (STRUCTURE)

```
HPCA_GRP_01_ASSGN
├── analysis.py
├── blocked-matmul
├── blocked-matmul.c
├── config.py
├── m5out
│   ├── output_LQEntries_64_SQEntries_64_lld_size_32kB_lld_size_8kB_l2_size_512kB_bp_type_BiModeBP_ROBEntries_192_numIQEntries_64
│   │   └── stats.txt
│   ├── output_LQEntries_64_SQEntries_64_lld_size_32kB_lld_size_8kB_l2_size_512kB_bp_type_TournamentBP_ROBEntries_192_numIQEntries_64
│   │   └── stats.txt
│   ├── output_LQEntries_64_SQEntries_64_lld_size_64kB_lld_size_16kB_l2_size_256kB_bp_type_BiModeBP_ROBEntries_192_numIQEntries_64
│   │   └── stats.txt
│   ├── output_LQEntries_64_SQEntries_64_lld_size_64kB_lld_size_16kB_l2_size_256kB_bp_type_TournamentBP_ROBEntries_192_numIQEntries_64
│   │   └── stats.txt
│   ├── output_LQEntries_64_SQEntries_64_lld_size_64kB_lld_size_16kB_l2_size_512kB_bp_type_BiModeBP_ROBEntries_192_numIQEntries_64
│   │   └── stats.txt
│   ├── output_LQEntries_64_SQEntries_64_lld_size_64kB_lld_size_16kB_l2_size_512kB_bp_type_TournamentBP_ROBEntries_192_numIQEntries_64
│   │   └── stats.txt
│   ├── output_LQEntries_64_SQEntries_64_lld_size_64kB_lld_size_8kB_l2_size_256kB_bp_type_BiModeBP_ROBEntries_192_numIQEntries_64
│   │   └── stats.txt
│   ├── output_LQEntries_64_SQEntries_64_lld_size_64kB_lld_size_8kB_l2_size_256kB_bp_type_TournamentBP_ROBEntries_192_numIQEntries_64
│   │   └── stats.txt
│   ├── output_LQEntries_64_SQEntries_64_lld_size_64kB_lld_size_8kB_l2_size_512kB_bp_type_BiModeBP_ROBEntries_192_numIQEntries_64
│   │   └── stats.txt
│   └── output_LQEntries_64_SQEntries_64_lld_size_64kB_lld_size_8kB_l2_size_512kB_bp_type_TournamentBP_ROBEntries_192_numIQEntries_64
│       └── stats.txt
├── options.py
├── Report.pdf
├── requirements.txt
└── run.py
```

STEPS TO RUN

Follow the following steps to run the program:

1. Copy the HPCA_GRP_01_ASSGN inside the ~/gem5/configs folder.
2. Open the terminal inside the HPCA_GRP_01_ASSGN folder and run
`python3 run.py`
3. The output can be found inside the
~/gem5/configs/HPCA_GRP_01_ASSGN/outputs folder.
4. Now install seaborn, which is a library used for plotting, (this can be done inside a virtual environment or directly) by executing `pip install`

seaborn'.

5. To generate all the statistics and analyse the outputs, execute '**python3 analysis.py**'. This generates a directory **m5out/**, which will contain stats for best 10 configuration, **best_10_config.csv** that contains parameter combinations of best 10 configurations and **plots/** where all the plots will be present.

ANALYSIS

Total number of configuration combinations: $2^8 = 256$

The configurations were generated using the following variable parameters.

- **LQEntries:** 32, 64
- **SQEntries:** 32, 64
- **l1d size:** 32kB, 64kB
- **l1i size:** 8kB, 16kB
- **l2 size:** 256kB, 512kB
- **bp type:** TournamentBP, BiModeBP
- **ROBEntries:** 128, 192
- **numIQEntries:** 16, 64

Benchmark Program: blocked-matmul.c

The top 10 configuration combinations based with respect to CPI is given in the table below:

LQEntries	SQEntries	l1d_size	l1i_size	l2_size	bp_type	ROBEntries	numIQEntries	cpi
64	64	64kB	16kB	512kB	BiModeBP	192	64	0.443066
64	64	64kB	16kB	512kB	TournamentBP	192	64	0.443066
64	64	64kB	8kB	512kB	TournamentBP	192	64	0.443139
64	64	64kB	8kB	512kB	BiModeBP	192	64	0.443139

64	64	64kB	16kB	256kB	TournamentBP	192	64	0.443178
64	64	64kB	16kB	256kB	BiModeBP	192	64	0.443178
64	64	64kB	8kB	256kB	BiModeBP	192	64	0.44324
64	64	64kB	8kB	256kB	TournamentBP	192	64	0.44324
64	64	32kB	8kB	512kB	TournamentBP	192	64	0.444871
64	64	32kB	8kB	512kB	BiModeBP	192	64	0.444871

Analysis of best configurations:

The above configurations work best for the given benchmark program for the following reasons:

1. The program divides the matrices into 2X2 submatrices and performs block wise multiplication. Since it needs to access blocks of contiguous memory multiple times cache size plays a crucial role in the performance as evident in the above configuration combinations, CPI is reducing gradually with the increase of the size of both the L1 and L2 caches.
2. The TournamentBP and the BiModeBP appear an equal number of times in the top 10 configurations. The BiModeBP is a better choice for our benchmark program since the for loops used are mostly simple loops with the pattern T*NT which can be effectively predicted with a simple BiModeBP.
3. A larger ROB size helps the CPU to issue and execute more instructions in-flight increasing performance. The same can be observed in our case as well.
4. Since our program executes a lot of load and store instructions a larger load queue and store queue size help boost the performance as evident in the above configurations.
5. A larger Instruction Queue size lets the CPU prefetch more instructions. Therefore the CPU never has to wait for the IQ to get filled to issue the next instruction. This can be seen in our case of the Benchmark program as well.

PLOTS

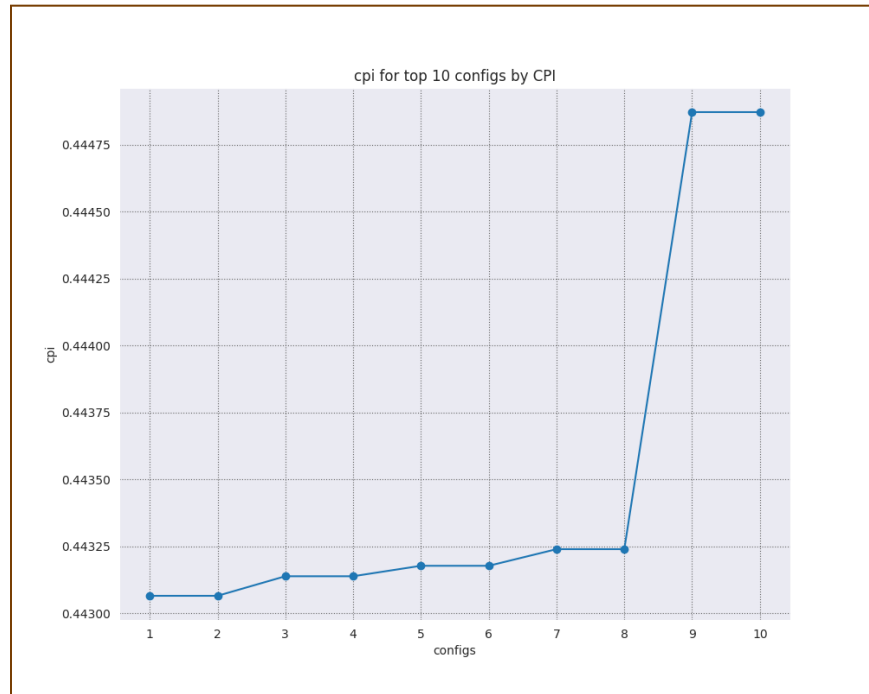


Figure 1: CPI for top 10 configurations




Figure 2: Mispredicted branches detected during execution




Figure 3: Number of branches that were predicted not taken incorrectly




Figure 4: Number of branches that were predicted taken incorrectly




Figure 5: Instructions Per Cycle (IPC)




Figure 6: Number of BTB hit percentage




Figure 7.a: Number of overall miss cycles - icache




Figure 7.b: Number of overall miss cycles - dcache




Figure 7.c: Number of overall miss cycles - l2cache




Figure 8.a: Overall miss rate - icache




Figure 8.b: Overall miss rate - dcache




Figure 8.c: Overall miss rate - l2cache




Figure 9.a: Overall average miss latency - icache




Figure 9.b: Overall average miss latency - dcache




Figure 9.c: Overall average miss latency - l2cache




Figure 10.a: The number of ROB accesses - read




Figure 10.b: The number of ROB accesses - write




Figure 11: Number of times the LSQ has become full, causing a stall




Figure 12: Number of loads that had data forwarded from stores




Figure 13: Number of times access to memory failed due to the cache being blocked